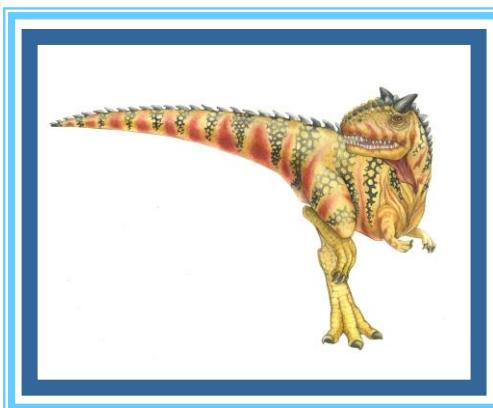
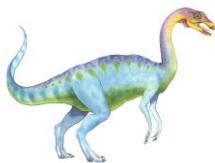


# Chapter 5: Synchronization





# Chapter 5: Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Semaphores
- Classic Problems of Synchronization
- Synchronization Examples





# Objectives

- To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
- To present both software and hardware solutions of the critical-section problem





# Background

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. Initially, **count is set to 0**. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

increment : เพิ่มค่าขึ้น  
decrement: ลดค่าลง





# Producer

```
while (true) {  
  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```





# Consumer

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item in nextConsumed  
}
```





# Race Condition

Recall that count is a shared variable

- count++ could be implemented as statement of the producer.

```
register1 = count          p1  
register1 = register1 + 1   p2  
count = register1          p3
```

- count-- could be implemented as statement of the consumer.

```
register2 = count          c1  
register2 = register2 - 1   c2  
count = register2          c3
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute register1 = count	{register1 = 5}
S1: producer execute register1 = register1 + 1	{register1 = 6}
S2: consumer execute register2 = count	{register2 = 5}
S3: consumer execute register2 = register2 - 1	{register2 = 4}
S4: producer execute count = register1	{count = 6 }
S5: consumer execute count = register2	{count = 4}

interleaving: การแทรกสลับการทำงานของชุดคำสั่ง





# Solution to Critical-Section Problem

A Solution must satisfy the following 3 requirements

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections (การห้ามอยู่พร้อมกัน)
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely (มีความก้าวหน้า) เวลาทั้งที่รู้ว่าจะต้องยกให้ในที่สุด
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted (รอค่อยอย่างมีขอบเขต)
  - Assume that each process executes at a nonzero speed
  - No assumption concerning relative speed of the  $N$  processes

Each process may consume different amount of time in execution of its critical section

Each process has a segment of code, called a critical

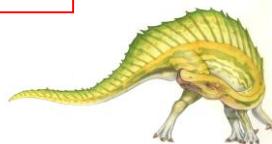
Critical section: เขตวิถีๆ คือพื้นที่ที่ process แต่ละตัวสามารถทำการปรับปรุงเปลี่ยนแปลงค่าตัวแปรต่างๆ ของ process โดยไม่มี process อื่นเข้ามาเกี่ยวข้องในพื้นที่นี้

exist: ยังปรากฏอยู่, คงอยู่

indefinitely : ไม่แน่นอน

postponed: ปฏิเสธ

granted: ได้รับการอนุญาตแล้ว





# Peterson's Solution

$P_i$  &  $P_j$

- Two process solution (เป็นวิธีที่ใช้กับ 2 process)
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:
  - int turn;
  - Boolean flag[2] use flag[0] & flag[1] to represent flag[i] & flag[j]
- The variable turn indicates whose turn it is to enter the critical section.  
= execute
- The flag array is used to indicate if a process is ready to enter the critical section. flag[i] = true implies that process  $P_i$  is ready!

1. เสียงต่ำในช่วงเวลาปกติ  
2. เสียงต่ำในช่วงเวลาผิดปกติ } 3 ตัวอย่าง  
3. เสียงต่ำในช่วงเวลาผิดปกติ }





# Algorithm for Process $P_i$

Initialization  
boolean flag[ $i$ ] - flag[ $i$ ] = flag[ $j$ ]

do {

```
flag[ $i$ ] = TRUE;  
turn =  $j$ ;  
while (flag[ $j$ ] && turn ==  $j$ );
```

critical section

```
flag[ $i$ ] = FALSE;
```

remainder section

} while (TRUE);

How can this Peterson's critical

มุ่งประดิษฐ์ให้ช่วงเวลาหนึ่งมี 2 process เท่านั้น  
solution meet 3 requirement (p.s.s) of a solution to the

See my note

$P_i$  does nothing and wait for  $P_j$  complete  
กระบวนการ  $P_i$  ไม่รบกวน (loop)

$i$ : current process  
 $j$ : other process





# Synchronization Hardware

- Many systems provide hardware support for critical section code
- Uniprocessors – could disable interrupts
  - Currently running code would execute without preemption
  - Generally too inefficient on multiprocessor systems
    - ▶ Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
  - ▶ Atomic = non-interruptable
  - Either test memory word and set value
  - Or swap contents of two memory words

Special H/W instruction that  
allow

Uniprocessor : โปรเซสเซอร์ตัวเดียว





# Solution to Critical-section Problem Using Locks

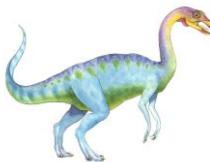
The simplest of higher-level S/W tool to solve the critical-section-problem is the mutex lock.

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

ต้องการล็อก

ปลดล็อก





# Semaphore

สำคัญ - มีผลลัพธ์ที่ต้องการ

A more robust tool that can behave similarly to a mutex lock (p.5.14)

- Synchronization tool that does not require busy waiting (ไม่ต้องการการรออยู่ที่มาก)
- Semaphore  $S$  – integer variable
- Two standard operations modify  $S$ : `wait()` and `signal()` ทดสอบ จงตัวตน การเพิ่มแบบเจ้าก็จะมาก
- Originally called `P()` and `V()` → to test      → to increment
- Less complicated  $P = \text{proberen}$        $V = \text{verhogen}$
- Can only be accessed via two indivisible (atomic) operations มี 2 การดำเนินการที่ทำกับSemaphore  
 $(S).wait \& signal(S)$

- **wait ( $S$ ) {**

```
    while  $S <= 0$ 
```

```
        ; // no-op
```

```
         $S--;$ 
```

```
}
```

Initialization  
 $S \geq 1$

Do nothing since the critical section is locked

}

Wait to test.

- **signal ( $S$ ) {**

```
     $S++;$  增加值
```

```
}
```

Signal to increment.





# Semaphore as General Synchronization Tool

- Counting semaphore – ~~integer~~ value can range over an unrestricted domain
- Binary semaphore – ~~integer~~ value can range only between 0 and 1; can be simpler to implement       $s \in [0,1]$  only when
- Also known as mutex locks
- Can implement a counting semaphore  $S$  as a binary semaphore
- Provides mutual exclusion

```
Semaphore mutex; // initialized to 1
do {
    wait (mutex);
    // Critical Section
    signal (mutex);
    // remainder section
} while (TRUE);
```

**wait (mutex):**

**while mutex≤ 0 do no-op;**  
**mutex--;**

wait to test      ~~soñariong n̄ m̄ 0 n̄ 1~~

**signal (mutex):**

**mutex++;**

signal to increment





# Semaphore Implementation

ໄຟລະກາກົມກວມປັບປຸງ

- Must guarantee that no two processes can execute wait () and signal () on the same semaphore at the same time since both operation update a semaphore(s.)
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
  - Could now have busy waiting in critical section implementation
    - ▶ But implementation code is short
    - ▶ Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution.

ໂຄສະນີດີຕໍ່ມັນນີ້ຈະ busy waiting ແກ້ໄຂ

ເຖິງການອາໄຫດ້ແລ້ວ  
ຮອດໃຈກຳນົດເຫັນ  
Yes

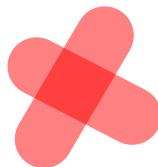




# Semaphore Implementation with no Busy waiting

ເຫດວ່າງານ

- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - block – place the process invoking the operation on the appropriate waiting queue. (ໃຫ້ເຂົ້າໄປອິນຄົວຍັ້ງໄຟ່ກໍານົດ)
  - wakeup – remove one of processes in the waiting queue and place it in the ready queue. (ນຳອອກຈາກຄົວເພື່ອຮອດກໍານົດ)



ຖົງນົມ  
2/9/66





# Semaphore Implementation with no Busy waiting (Cont.)

Initialization  
 $s = 1$

- Implementation of wait:

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

ค่า value อาจติดลบได้แสดงให้เห็นว่ามี process รออย semaphore

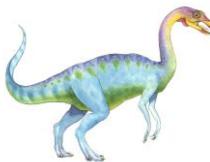
A linked list is used as a waiting queue of semaphore.

- Implementation of signal:

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Semaphor 5.1 52. ဝါယာမေးခွာ

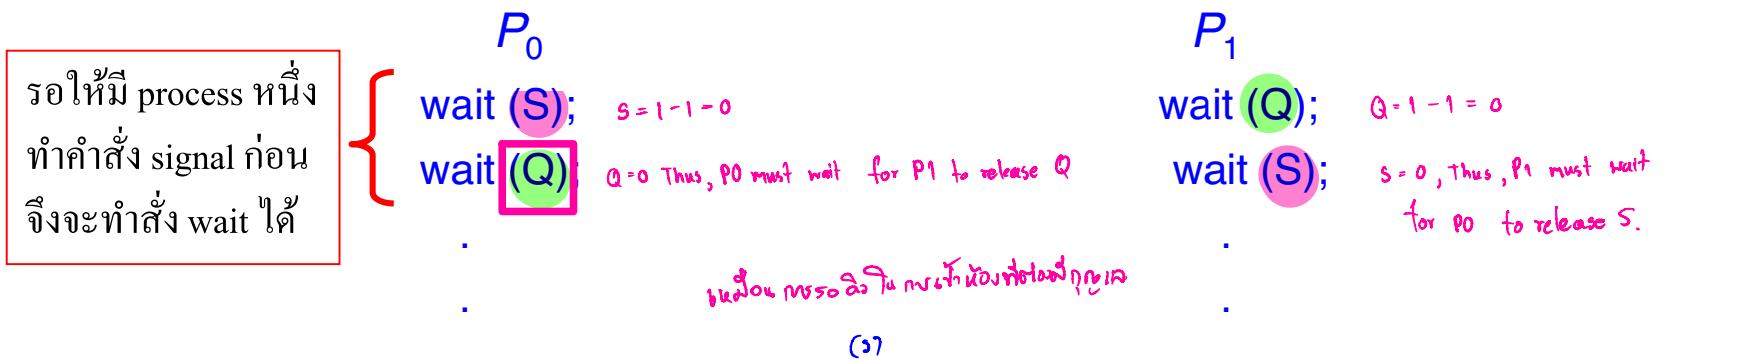




# Deadlock and Starvation

การใช้ **semaphore** อาจทำให้เกิดเหตุการณ์ขึ้นได้ ดังนี้

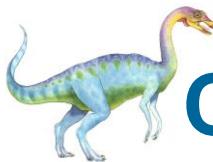
- Deadlock – two or more processes are **waiting indefinitely** for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1



- Starvation – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended
- Priority Inversion - Scheduling problem when lower-priority process holds a lock needed by higher-priority process

มีคนต้องรักษาอยู่  $\hookrightarrow$  คนต้อง Lower

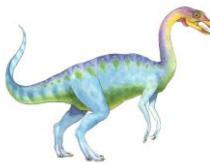




# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem





# Bounded-Buffer Problem

- $N$  buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N.

Producer & consumer





# Bounded Buffer Problem (Cont.)

wait &

- The structure of the **producer process**

do {

// produce an item in nextp

empty = 5 → 4 → 3 → 2 → 1 → 0

wait (empty);

wait (mutex);

// add the item to the buffer

mutex = 0 → 1  
signal (mutex);

signal (full);

} while (TRUE);

Initialization:

int N=5

semaphore

Empty --



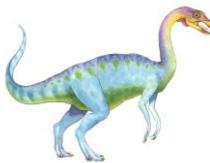
The 1st producer lock the execution of critical section (code of buffer updating)  
mutex = 1 → 0

Full ++



Full = 0 → 1 → 2 → 3 → 4 → 5.



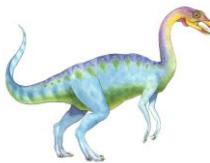


# Bounded Buffer Problem (Cont.)

## ■ The structure of the consumer process

```
do {      full = 5 → 4 → 3 → 2 → 1 → 0
    wait (full);
    mutex = 1 → 0
    wait (mutex);
    // remove an item from buffer to nextc
    mutex = 0 → 1
    signal (mutex);      The 1st consumer releases lock
    signal (empty);      Critical section
    empty = 0 → 1 → 2 → 3 → 4 → 5
    // consume the item in nextc
} while (TRUE);
```





# Readers-Writers Problem

- ใช้ข้อมูลร่วมกัน ผู้อ่านสามารถอ่าน (Reader) ข้อมูลร่วมกันได้หลาย ๆ คน
- ผู้เขียน (Writer) 1 คน สามารถเขียนข้อมูลได้ ณ ช่วงเวลาหนึ่ง โดยไม่มีผู้เขียนคนอื่นมาใช้ข้อมูลร่วม และห้ามผู้อ่านมาอ่านขณะที่เขียนอยู่

อาจทำให้เกิดปัญหา **Starvation** ได้ทั้ง สิ่งผู้เขียน และสิ่งผู้อ่าน

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write
- **Problem** – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time
- Shared Data
  - Data set *for reading*
  - Semaphore **mutex** initialized to 1
  - Semaphore **wrt** initialized to 1 *(or mutex)*
  - Integer **readcount** initialized to 0

ป้องกันตัวแปร **readcount** (ผู้อ่าน)

ป้องกันผู้เขียน





# Readers-Writers Problem (Cont.)

ឧបនៃរួម និងរួមក្នុងការសរស់សិក្សា

## ■ The structure of a writer process

```
lock the critical-section  
do {  
    wait (wrt) ;
```

- ធ្វើតាមការក្លាយទៅតាមការសរស់សិក្សាដែលត្រូវបានបញ្ជាក់  
wrt ដើម្បីធ្វើការសរស់សិក្សាដែលត្រូវបានបញ្ជាក់

```
// writing is performed
```

```
    signal (wrt) ;  
} while (TRUE);
```





# Readers-Writers Problem (Cont.)

## The structure of a reader process

do {

    mutex = 1 → 0

    wait (mutex) ;

    readcount ++ ;    readcount = 0 → 1 → 2 → 3 → 4 → ... → n  
    if (readcount == 1)  
        wait (wrt) ;  
    iron reader

    signal (mutex)    mutex = 0 → 1

เขียนโปรแกรม

หากผู้อ่านมีมากกว่า 1 คน ถ้า  
ผู้เขียนกำลังทำงานอยู่ ผู้อ่านคน  
ที่ 2 จะรออยู่ โดยการตรวจสอบ  
ตัวแปร mutex

Critical section #1 (block other reader from updating read count & wrt  
(between wait(mutex)... signal(mutex))

// reading is performed

Critical section #3 (block writes)  
(between wait(wrt)... signal(wrt))

    wait (mutex) ;    mutex = 1 → 0

    readcount -- ;

    if (readcount == 0)

        signal (wrt) ;

    signal (mutex) ;

readcount = m → m-1 → m-2 → ... → 2 → 1 → 0  
the last reader release the lock for writers

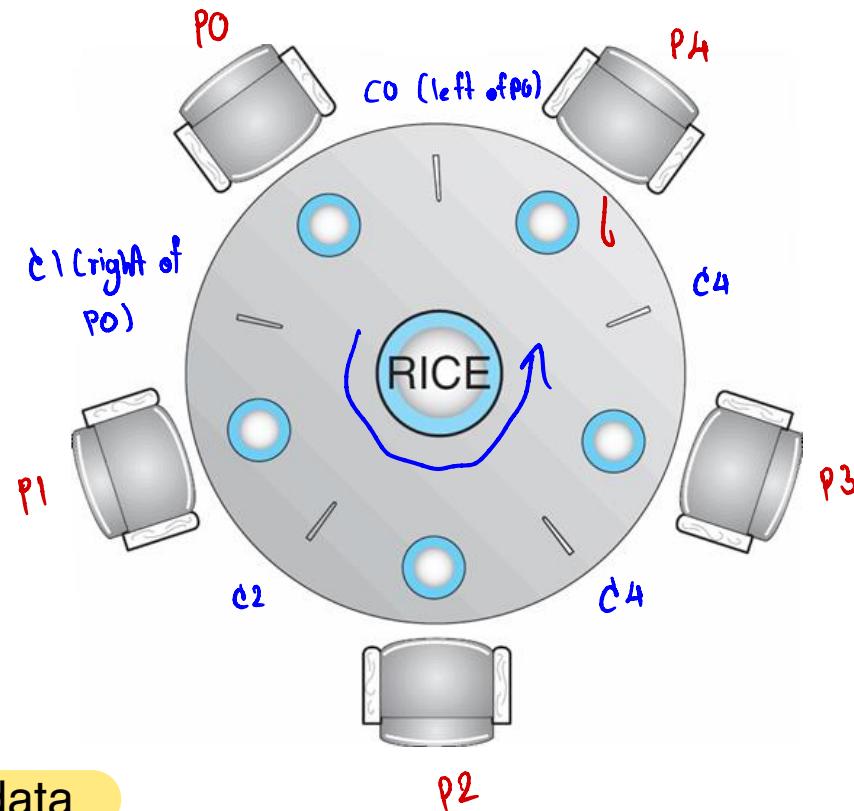
Critical section #2 (block other reader from  
updating readcount & wrt)  
(between wait(mutex)... signal(mutex))

} while (TRUE);





# Dining-Philosophers Problem

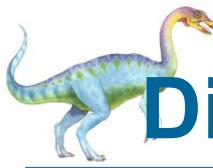


- Shared data

- Bowl of rice (data set)
- Semaphore chopstick [5] initialized to 1

c0 ... c4





# Dining-Philosophers Problem (Cont.)

## ■ The structure of Philosopher *i*:

```
do {
```

*i* in [0...4]

wait ( chopstick[*i*] ); lock Pi's left chopstick

wait ( chopStick[ (*i* + 1) % 5] ); lock Pi right chopstick

// eat

signal ( chopstick[*i*] );

signal (chopstick[ (*i* + 1) % 5] );

// think

Release Pi's right  
Chopstick

```
} while (TRUE);
```

หยິບຕະເກີຍນໃໝ່  
**operation Wait**

ວາງຕະເກີຍນໃໝ່  
**operation Signal**

ອາຈເກີດປ້ມງູຫາ  
**Deadlock** ໄດ້ ລາກ

ທຸກຄົນທີ່ໄວພຣື້ອມກັນ

ແລ້ວຫຍິບຕະເກີຍນຂ້າງໜ້າຍ

ເໜືອນກັນໜຳດ



# Dining-Philosophers Problem (Cont.)

- อาจเกิดปัญหา **Deadlock** ได้ หากทุกคนหิวพร้อมกัน

แล้วหยิบตะเกียงข้างซ้ายเหมือนกันหมด

วิธีแก้ไขเพื่อเลี่ยงการเกิด **Deadlock**

- \* มีนักประช蜃์นั่งโต๊ะ ได้ไม่เกิน 4

- \* กำหนดให้จะหยิบตะเกียงได้ตະเกียงด้านซ้ายและขวาต่อไปเรื่อยๆ (ขณะอยู่ใน

## Critical-Section

- \* ใช้การสลับกัน เช่น ให้คนเลขคี่หยิบซ้ายก่อน ข้างขวา และให้คนเลขคู่ หยิบขวา ก่อน ข้างซ้าย

\*\* อาจเกิดปัญหา **Starvation** ได้หากแก้ไขไม่รัดกุม \*\*





# Problems with Semaphores

- incorrect use of semaphore operations:

(การใช้ **operation** ของ **semaphore** ที่ไม่ถูกต้อง)

- signal (mutex) .... wait (mutex) ทำให้ไม่เกิดคุณสมบัติ Mutual exclusion
- wait (mutex) ... wait (mutex) ทำให้เกิดปัญหา Deadlock ได้ เพราะ ไม่มีใครปลดล็อก
- Omitting of wait (mutex) or signal (mutex) (or both)  
มีการละเลยการใช้ operation wait() หรือ signal() หรือทั้งคู่ จึงทำให้กลไกการทำงานของ semaphore ไม่เข้าจังหวะกัน





# Synchronization Examples

- Solaris
- Windows XP
- Linux
- Pthreads    *xI library*





# Solaris Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- Uses adaptive mutexes for efficiency when protecting data from short code segments  
See example on p.5.34
- Uses condition variables and readers-writers locks when longer sections of code need access to data  
For a long critical section to access the shared data
- Uses turnstiles to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock





# Windows XP Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems
  - Spinlock = a thread waits in a loop or spin unit the lock is available (released)*
- Uses **spinlocks** on multiprocessor systems
  - For efficiency, the kernel ensures that a thread will never be preempted while holding a spinlock*
- Also provides **dispatcher objects** which may act as either **mutexes** and **semaphores**
  - For thread synchronization, a dispatcher object may use diff mechanism including mutex lock,*
- Dispatcher objects may also provide **events**
  - An event acts much like a condition variable

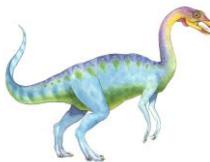




# Linux Synchronization

- Linux:
  - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
  - Version 2.6 and later, fully preemptive
  
- Linux provides:
  - semaphores
  - spin locks





# Pthreads Synchronization

Many system that implement Pthread also provide semaphore  
Semaphore are not part of the POSIX

- Pthreads API is OS-independent
- It provides:
  - mutex locks
  - condition variables
- Non-portable extensions include:
  - read-write locks
  - spin locks



# End of Chapter 5

