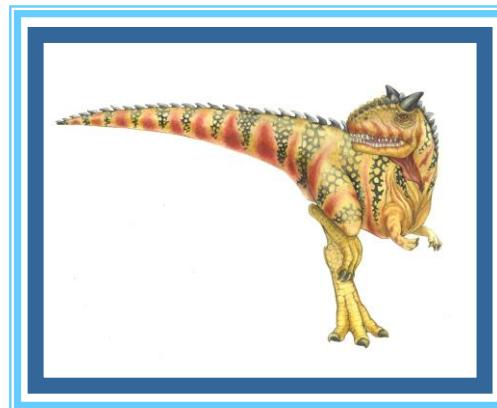
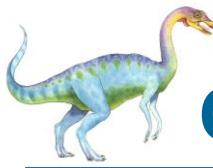


Chapter 8:

Virtual-Memory Management





Chapter 8: Virtual-Memory Management

- Background
- Demand Paging
- Page Replacement
- Allocation of Frames





Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames





Background

■

- Code needs to be in memory to execute, but entire program rarely used
 - Error code, unusual routines, large data structures
- Entire program code not needed at same time
- Consider ability to execute partially-loaded program
 - Program no longer constrained by limits of physical memory
 - Each program takes less memory while running -> more programs run at the same time
 - ▶ Increased CPU utilization and throughput with no increase in response time or turnaround time
 - Less I/O needed to load or swap programs into memory -> each user program runs faster





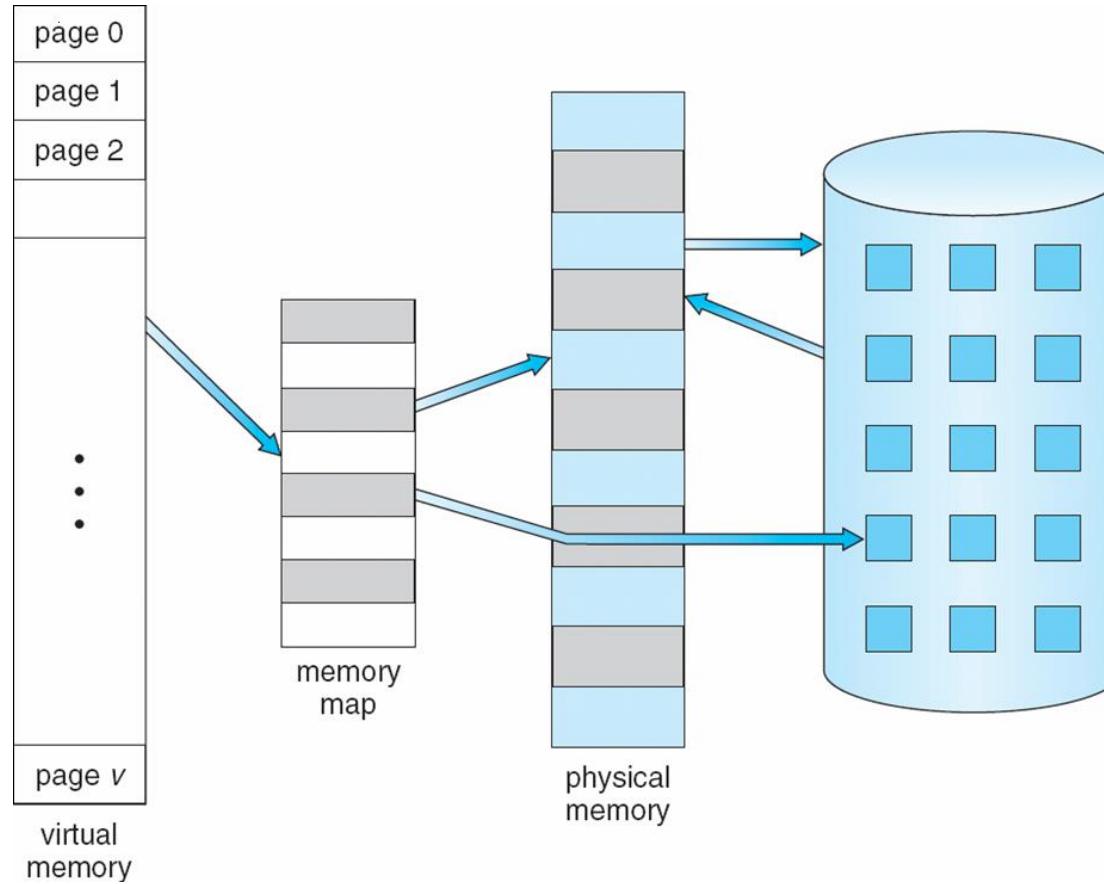
Background

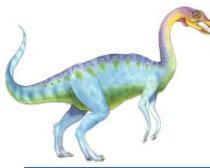
- **Virtual memory** – separation of user logical memory from physical memory.
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
 - More programs running concurrently
 - Less I/O needed to load or swap processes
- **Virtual address space** – logical view of how process is stored in memory
 - Usually start at address 0, contiguous addresses until end of space
 - Meanwhile, physical memory organized in page frames
 - MMU must map logical to physical
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation



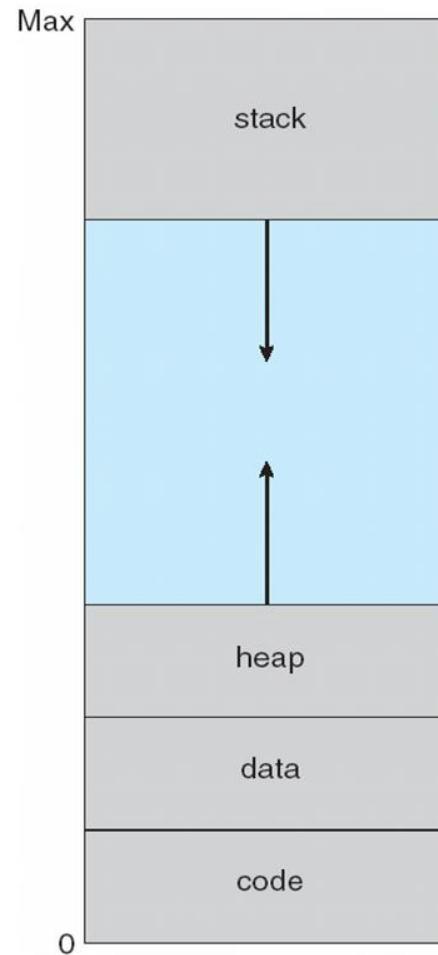


Virtual Memory That is Larger Than Physical Memory



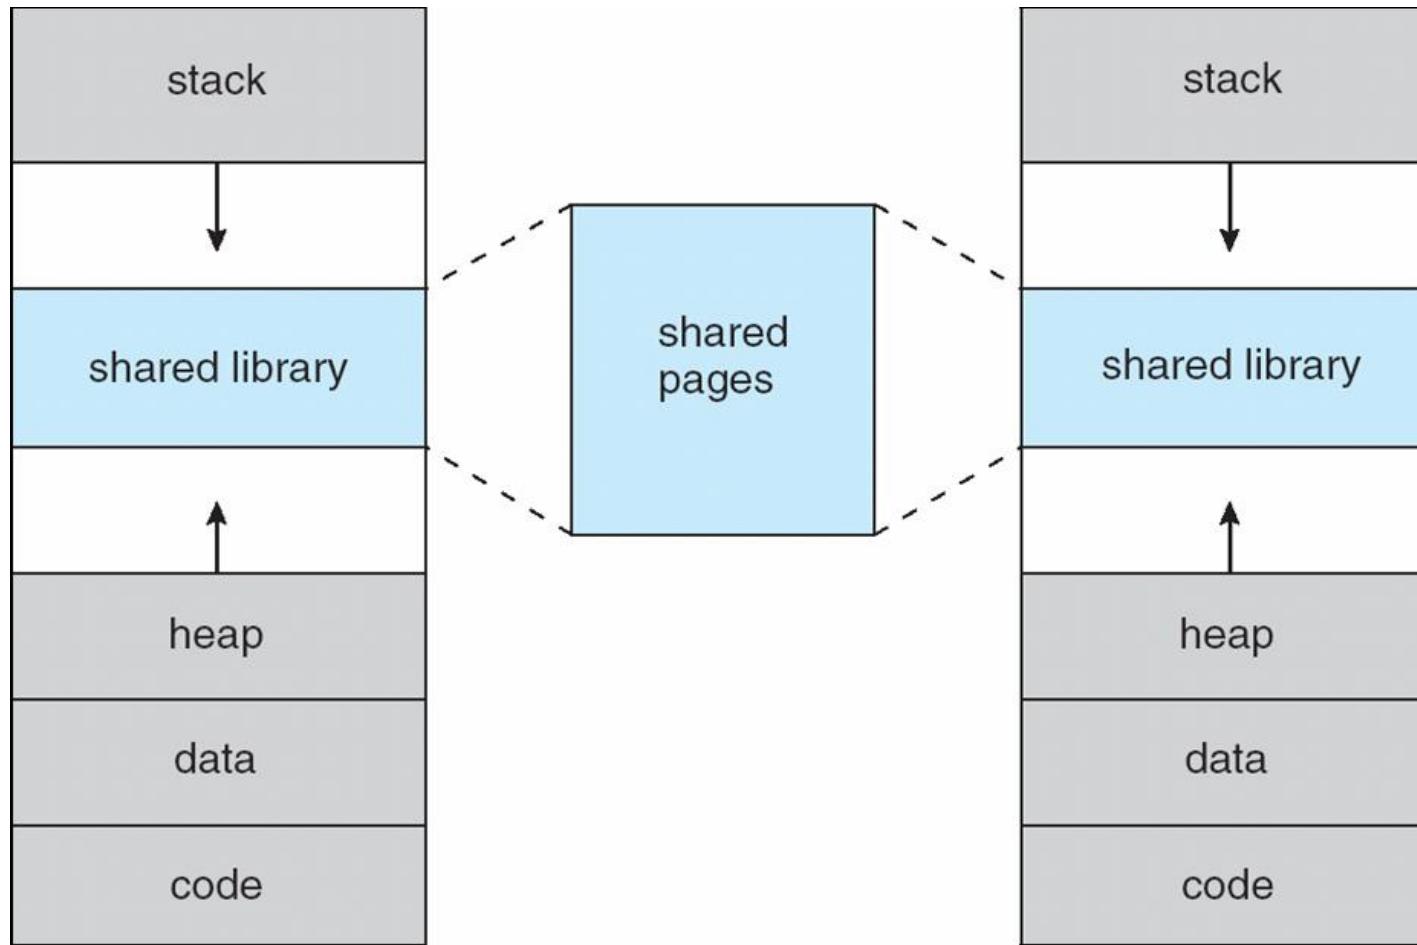


Virtual-address Space





Shared Library Using Virtual Memory





Demand Paging

(การจัดสรร **Paging** ตามความต้องการที่ร้องขอ)

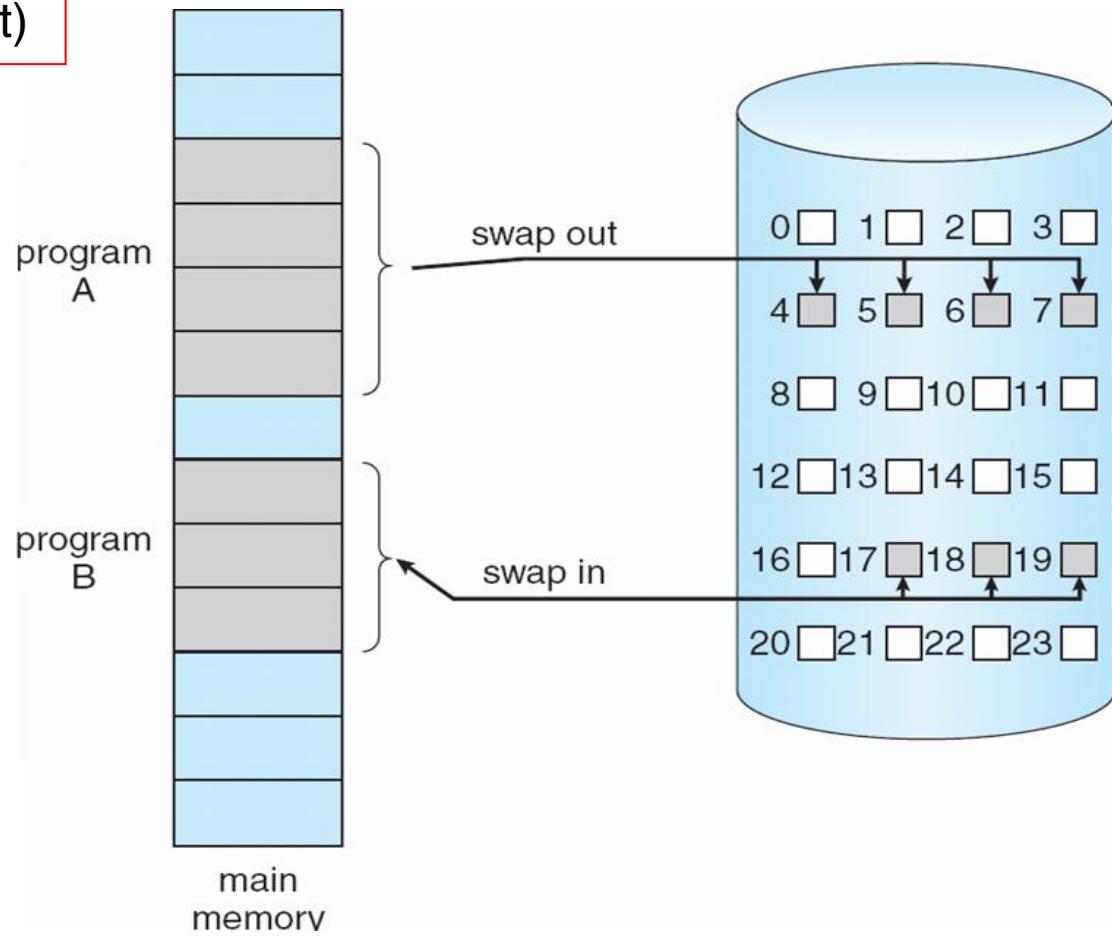
- Bring a page into memory **only when it is needed**
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed ⇒ reference to it (หากต้องการใช้ page ให้อ้างตำแหน่งลงใน page ที่ต้องการ)
 - invalid reference ⇒ abort (หากอ้างตำแหน่งไม่ถูกต้องให้ยกเลิก)
 - not-in-memory ⇒ bring to memory (หากอ้างแล้วไม่มีใน memory ให้นำเข้ามายังใน memory)
- **Lazy swapper** – never swaps a page into memory unless page will be needed
 - Swapper that deals with pages is a **pager**





Transfer of a Paged Memory to Contiguous Disk Space

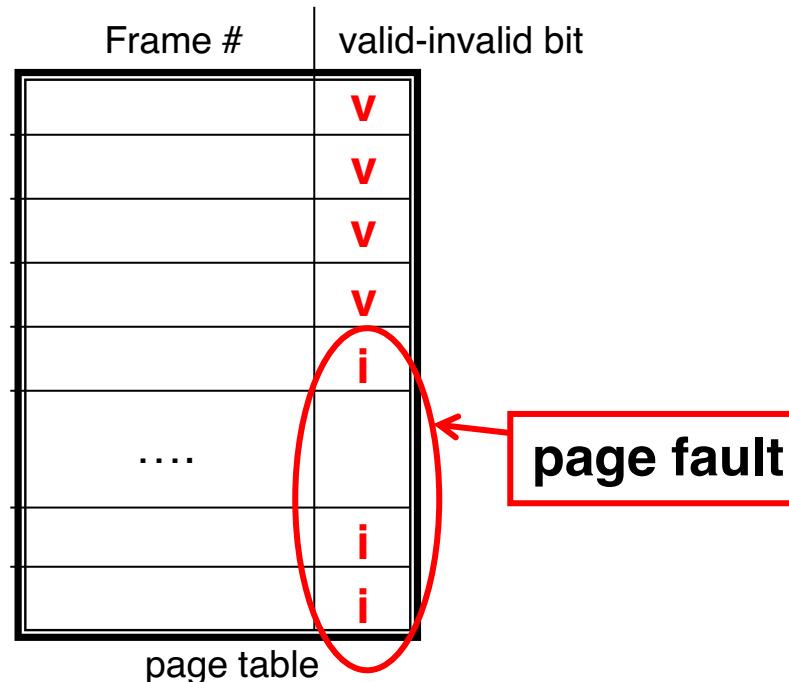
มี pager เป็นตัวจัด page
เข้า (swap in) หรือออก (swap out)





Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated (**v** ⇒ in-memory, **i** ⇒ not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:

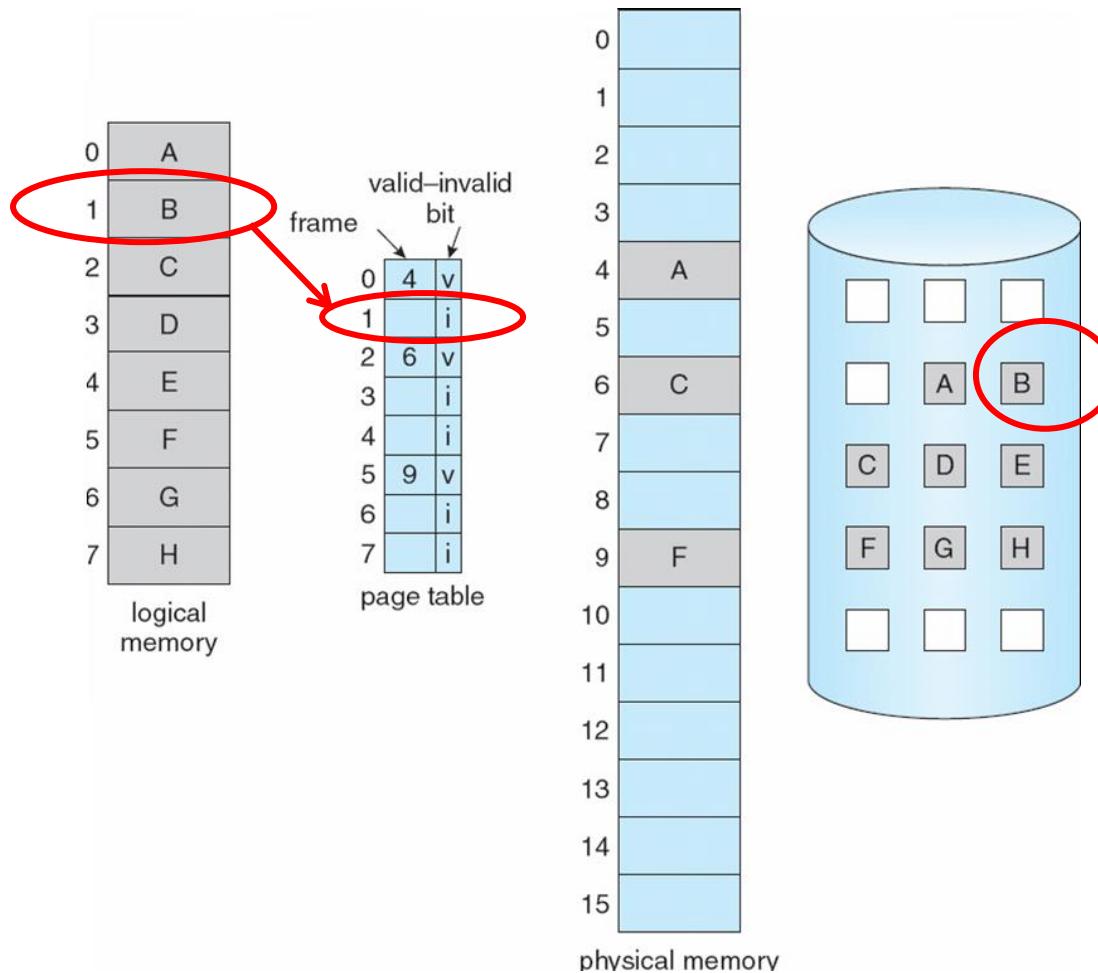


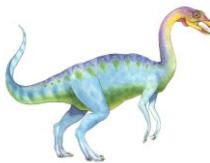
- During address translation, if valid–invalid bit in page table entry is **i** ⇒ page fault (การอ้างอิงพิกัดหน้าหรือไม่พบหน้าที่ต้องการ).





Page Table When Some Pages Are Not in Main Memory





Page Fault



เมื่อมีการอ้างอิงพิดหน้าหรือไม่พบหน้าที่ต้องการในหน่วยความจำหลัก (**page fault**) จะมีขั้นตอนดำเนินการดังนี้

- If there is a reference to a page, first reference to that page will trap to operating system:

page fault

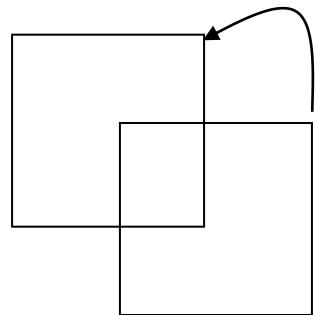
1. Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
2. Get empty frame
3. Swap page into frame
4. Reset tables
5. Set validation bit = **v**
6. Restart the instruction that caused the page fault





Page Fault (Cont.)

- Restart instruction (ทำต่อจากจุดที่ได้ทำมาแล้วล่าสุด หรือ ทำต่อจากจุดที่เกิด page fault ขึ้น)
 - block move

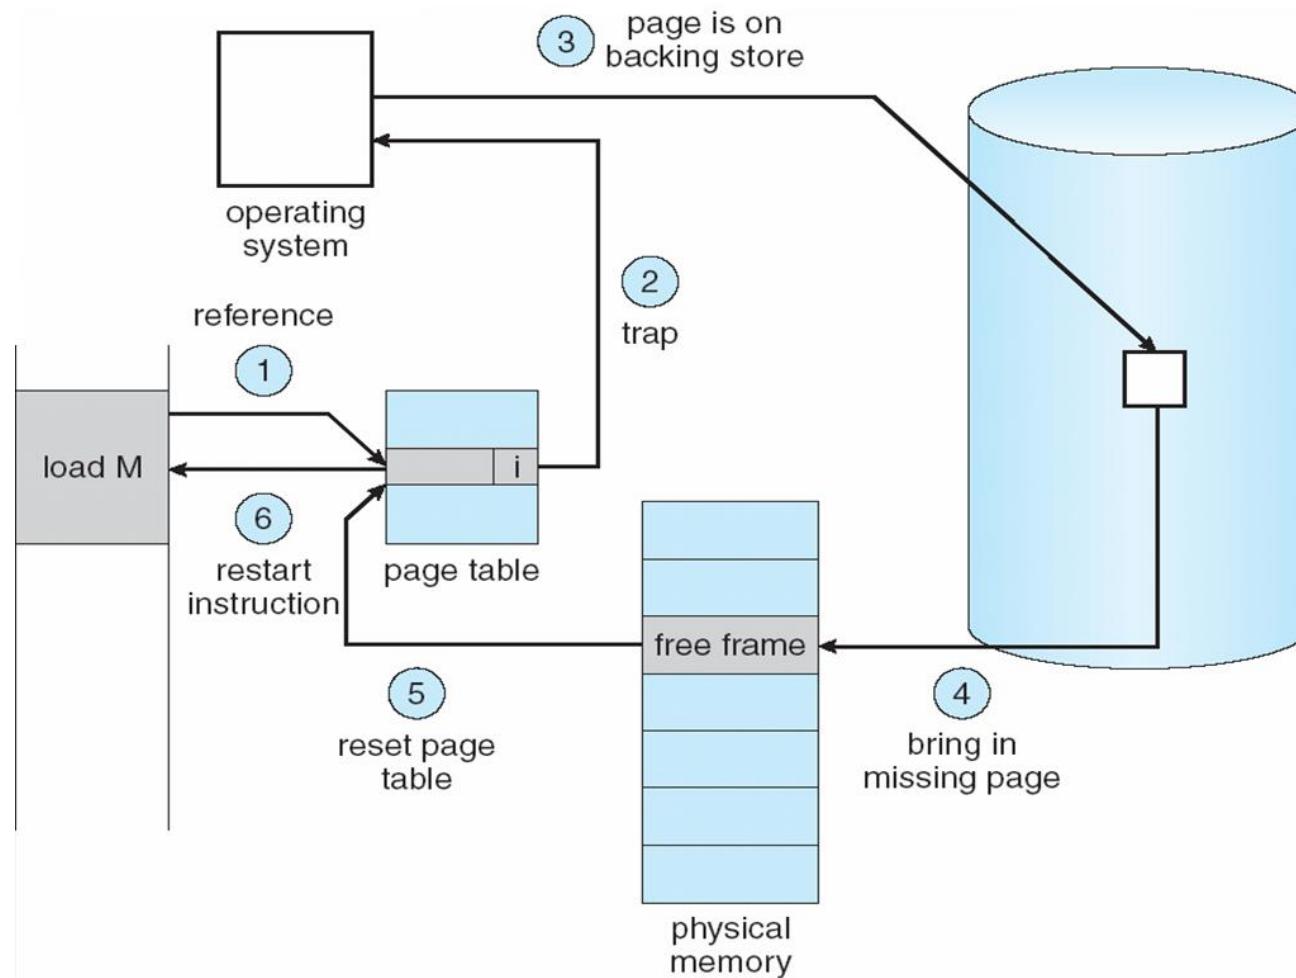


- auto increment/decrement location





Steps in Handling a Page Fault





What happens if there is no free frame?

- **Page replacement** – find some page in memory, but not really in use, swap it out
 - algorithm
 - performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

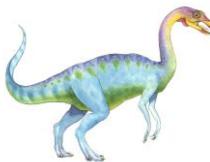




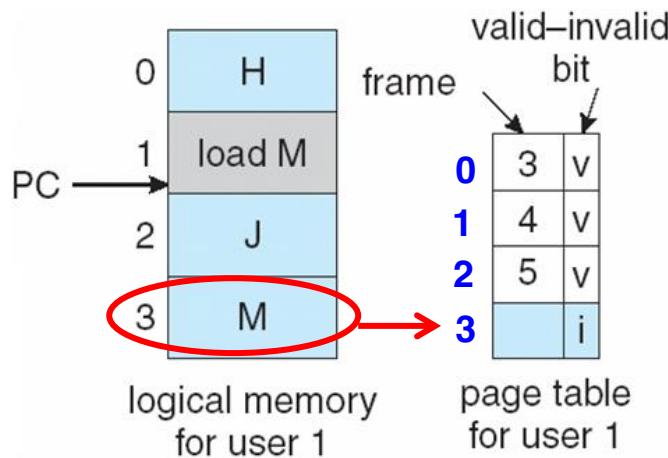
Page Replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- Use modify (dirty) bit to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory



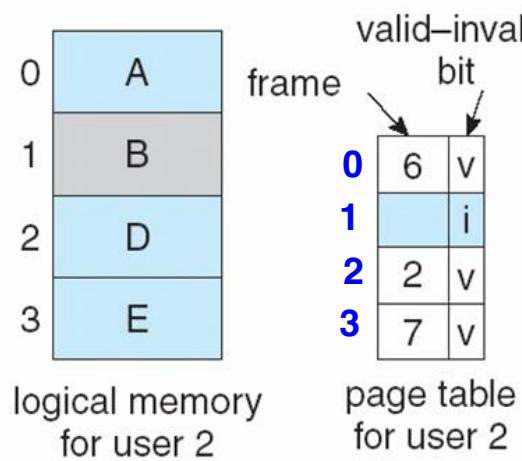
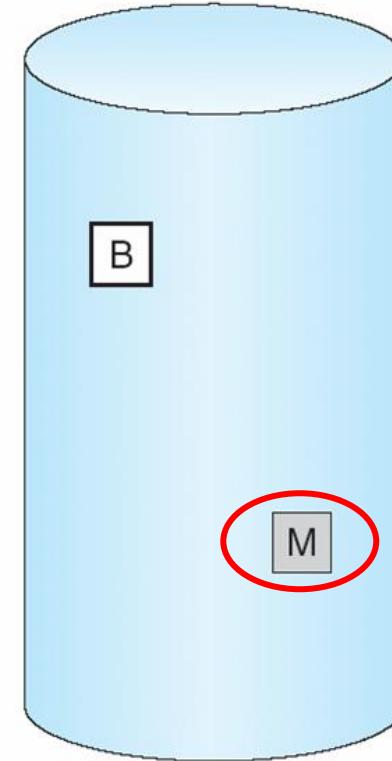


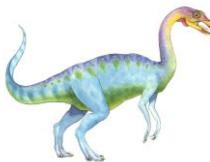
Need For Page Replacement



physical memory

0	monitor
1	D
2	H
3	load M
4	J
5	A
6	E
7	

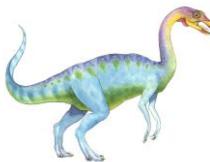




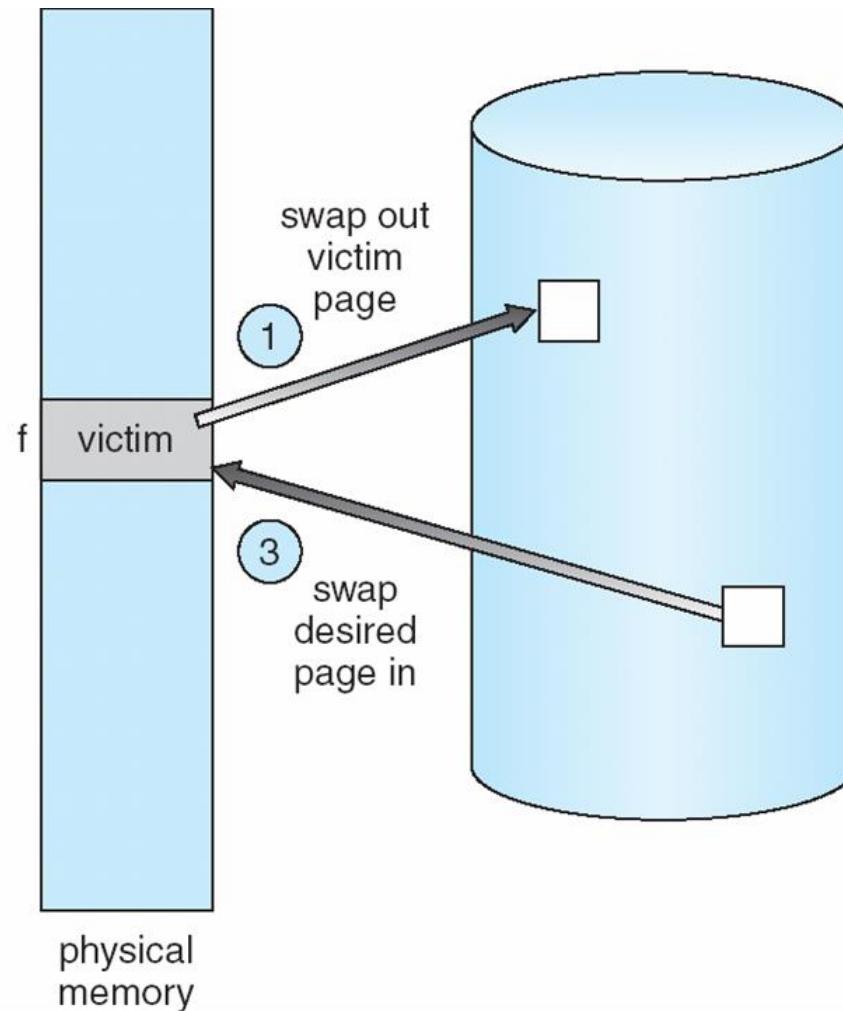
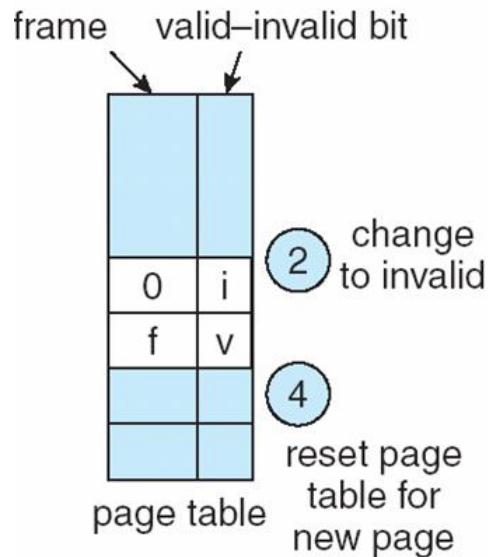
Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process





Page Replacement





Page Replacement Algorithms

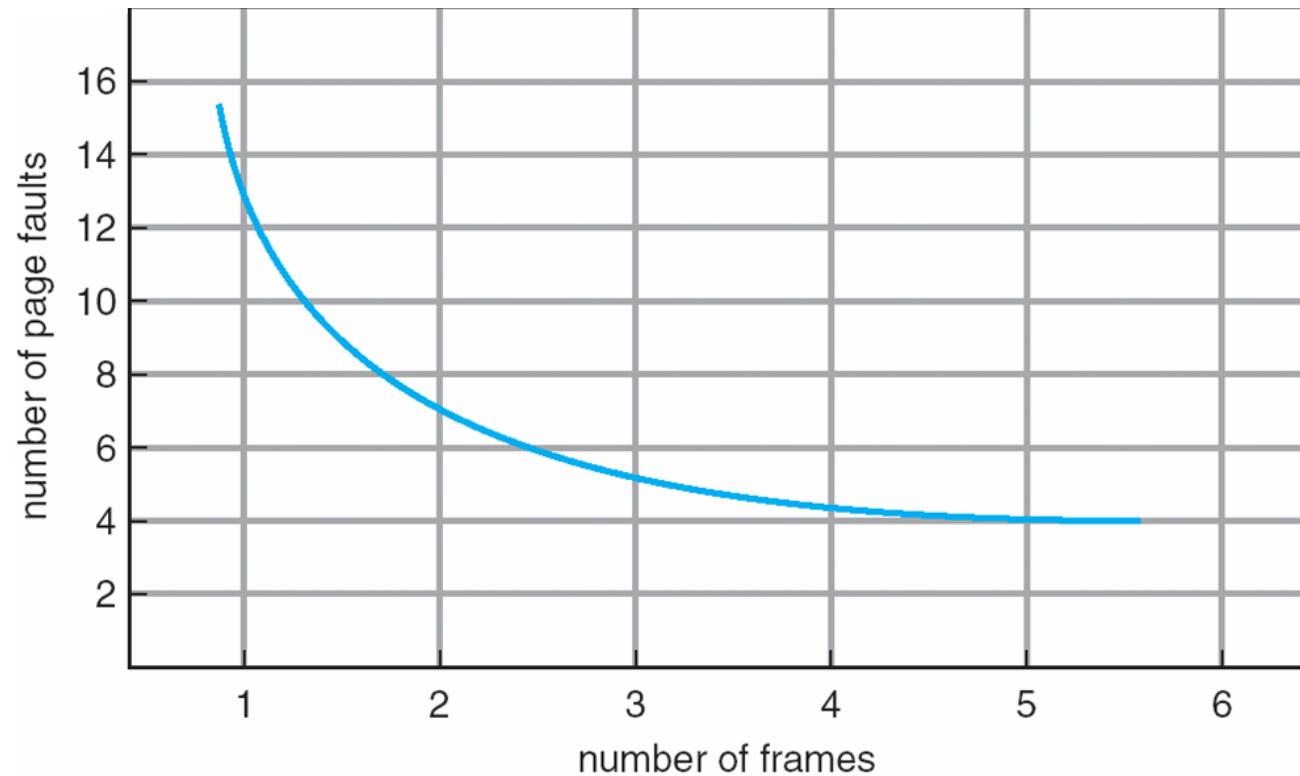
- Want lowest page-fault rate
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5





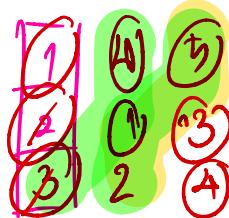
Graph of Page Faults Versus The Number of Frames





First-In-First-Out (FIFO) Algorithm

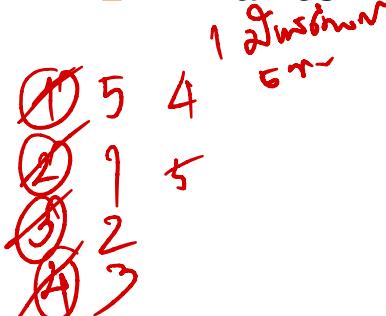
- Reference string: ~~1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5~~
- 3 frames (3 pages can be in memory at a time per process)



frame #	1	2	3	4	5
1	1				
2		2		1	3
3			3	2	4

9 page faults

- 4 frames



frame #	1	2	3	4	5	4
1	1					
2		2			1	5
3			3		2	2
4				4		3

10 page faults

see my note

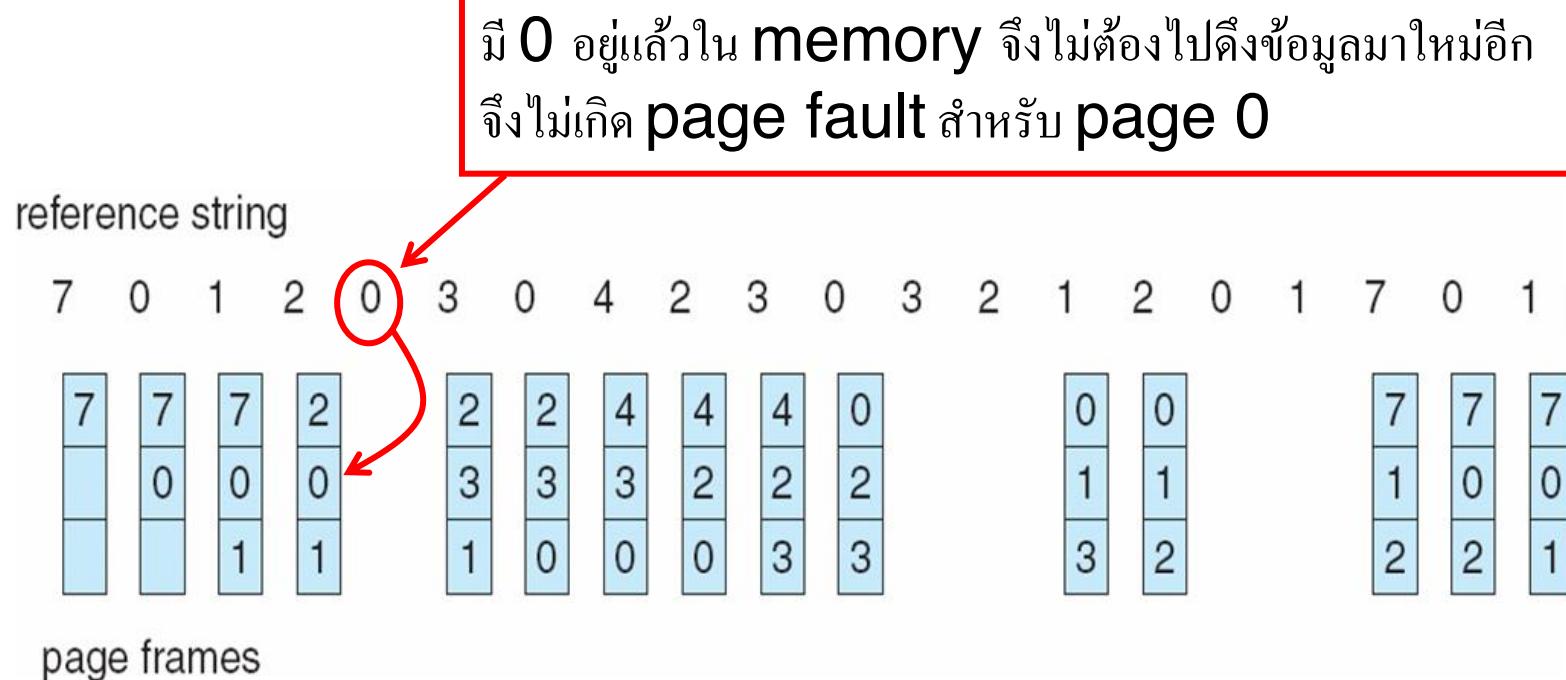
10 page

- Belady's Anomaly: more frames \Rightarrow more page faults





FIFO Page Replacement



เมื่อ physical memory 3 frames และใช้ reference string
ตามที่กำหนดให้ จะเกิด page fault ทั้งหมด 15 page faults

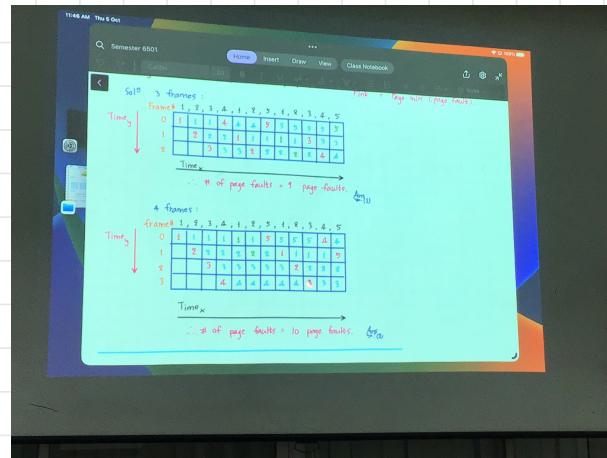
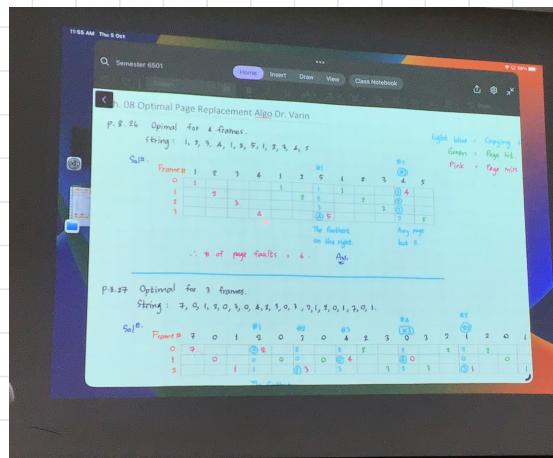
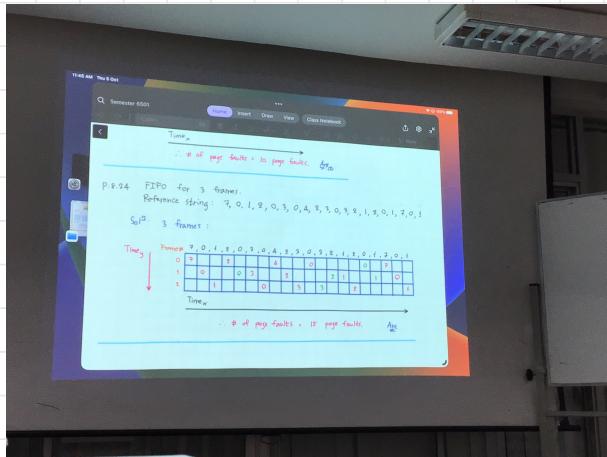
- ครั้งแรกไม่มี page ออย์ใน Physical Memory จะเกิด page fault



7.0 X 2 080 X 23.0 32.820 870 X

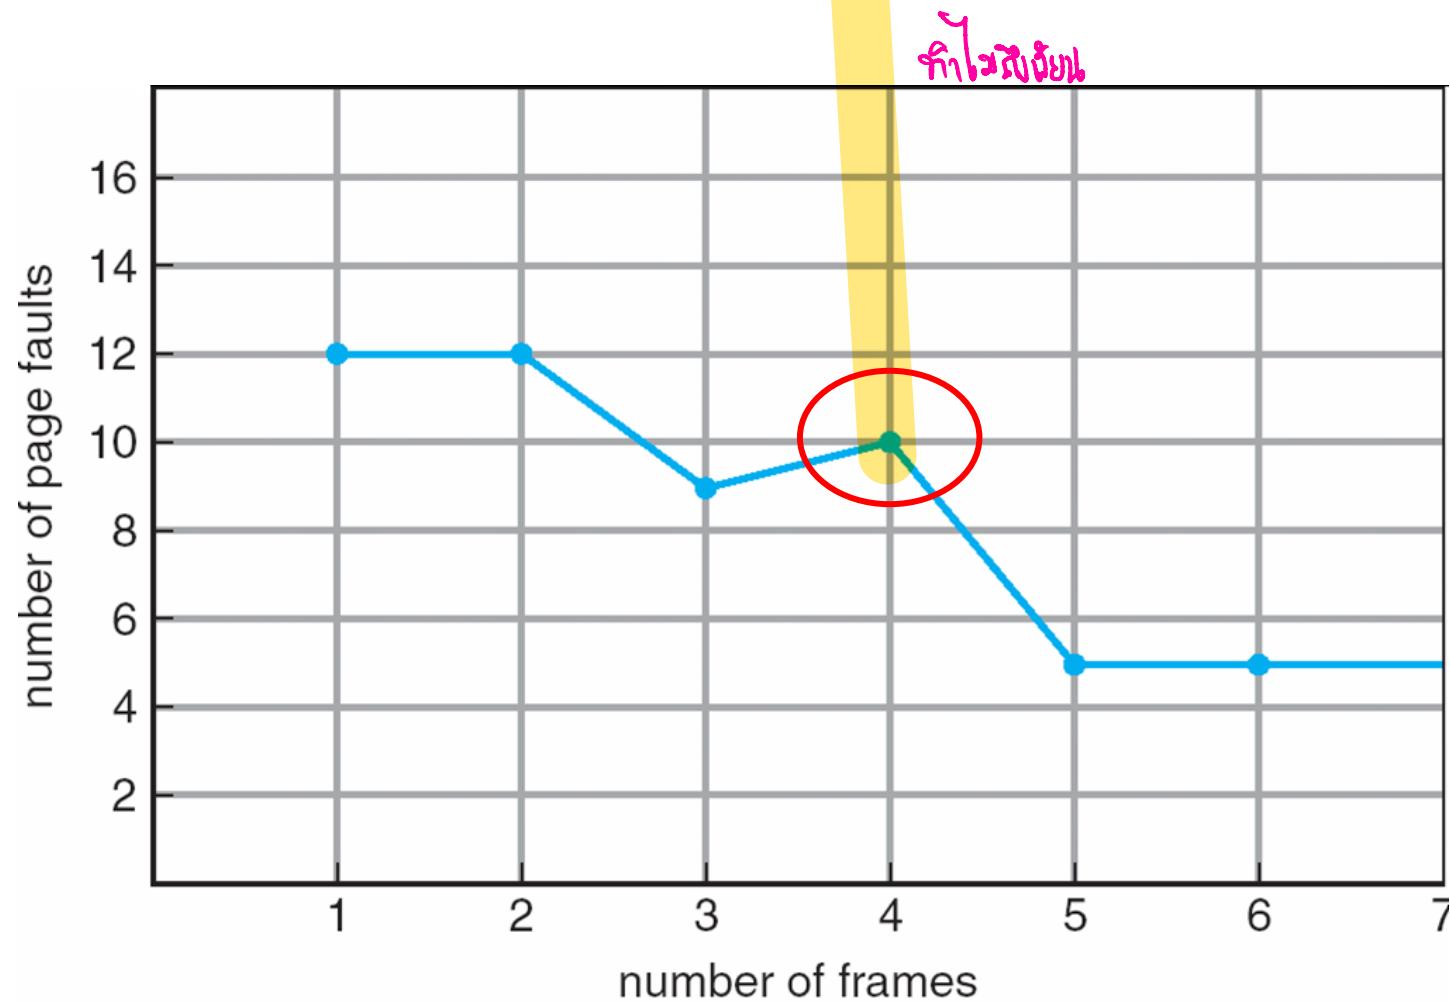
X	2	4	0	7
0	3	2	X	0
X	0	3	2	1

} 15 page faults



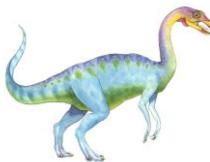


FIFO Illustrating Belady's Anomaly



Belady's Anomaly : เป็นเหตุการณ์ที่เมื่อมี Physical Memory เพิ่มขึ้น แต่จะเกิด page fault เพิ่มขึ้นด้วย (เป็นข้อยกเว้นสำหรับ FIFO Algorithm)





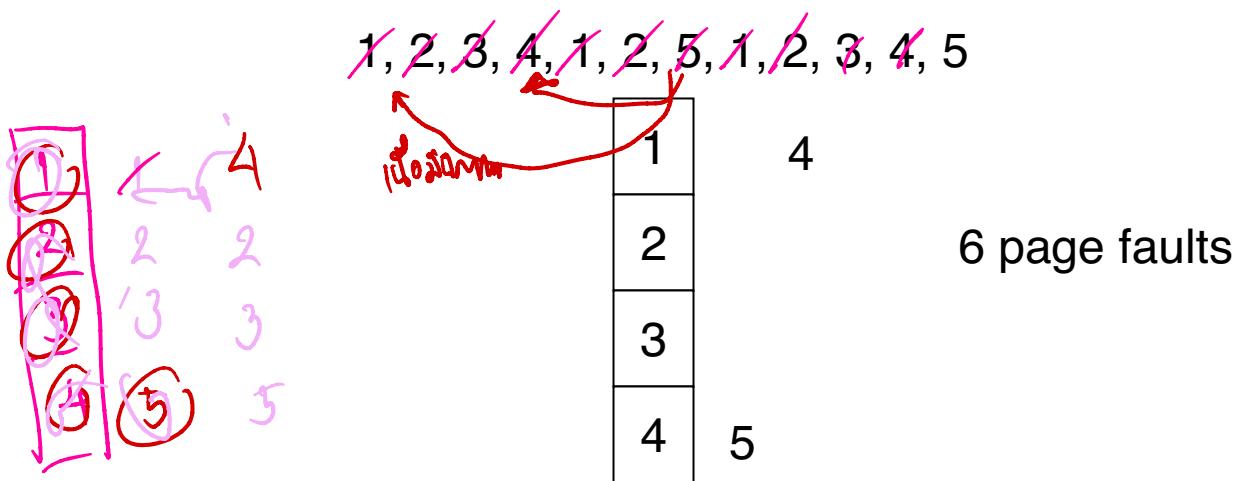
Optimal Algorithm

Victim = the farthest on the right

- Replace page that will not be used for longest period of time

(มองใน list ของ reference string ถัดไปในอนาคตว่า page ใดอีกหนานที่จะถูกใช้งาน จะโดน replace)

- 4 frames example



- How do you know this? ไม่ make sense เราไม่สามารถรู้อนาคตได้
- Used for measuring how well your algorithm performs

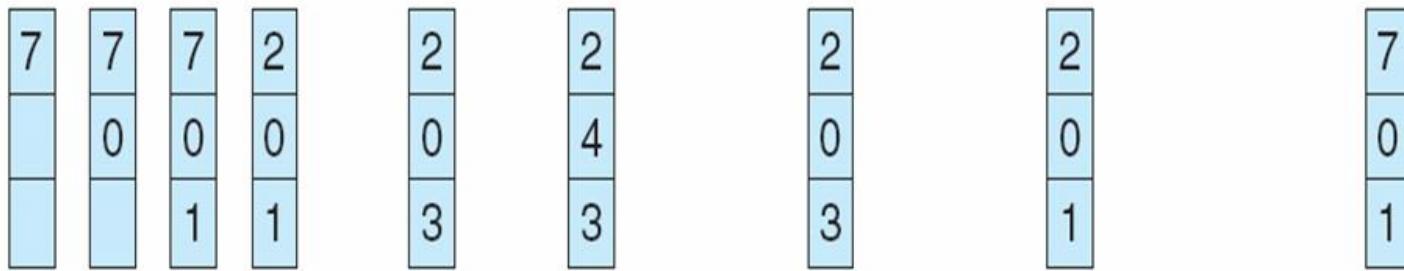




Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

เกิดกี่ page fault ???

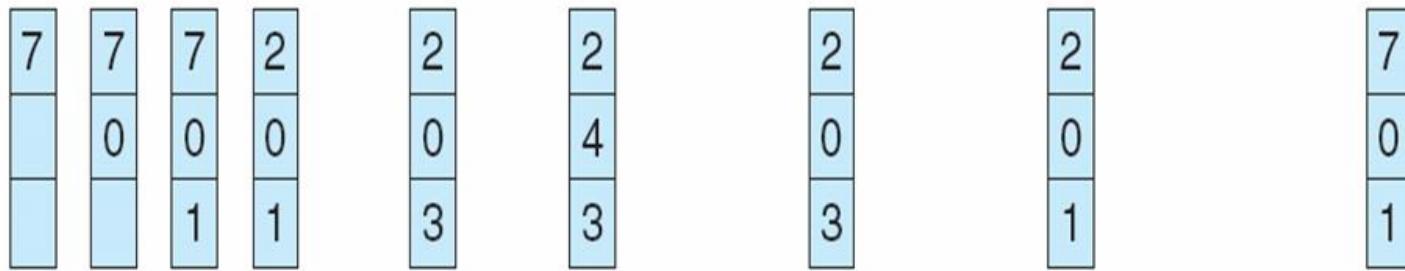




Optimal Page Replacement

reference string

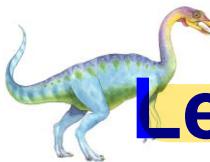
7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

เกิด 9 page fault

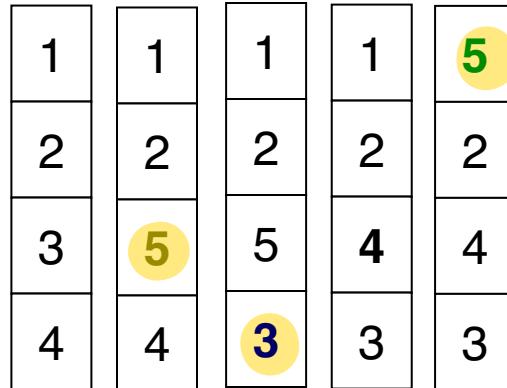




Least Recently Used (LRU) Algorithm

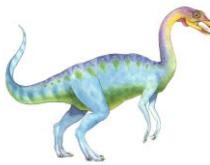
(มองย้อนกลับไปในอ็ดิตว่า page ได้ไม่ได้ถูกใช้งานนานที่สุดจะถูก replace)

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters to determine which are to change

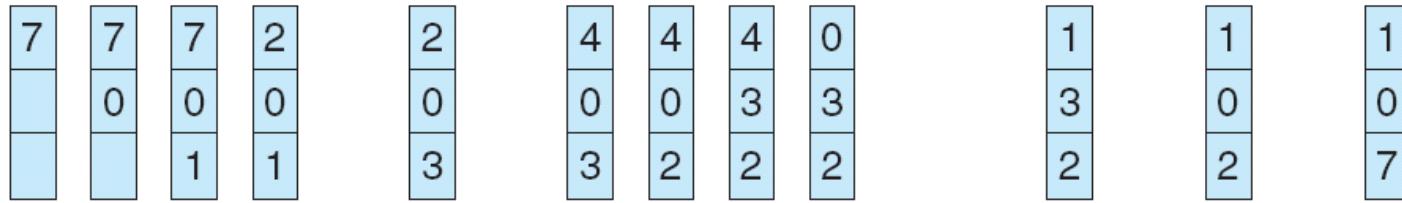




LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

72 30)

เกิดกี่ page fault ???

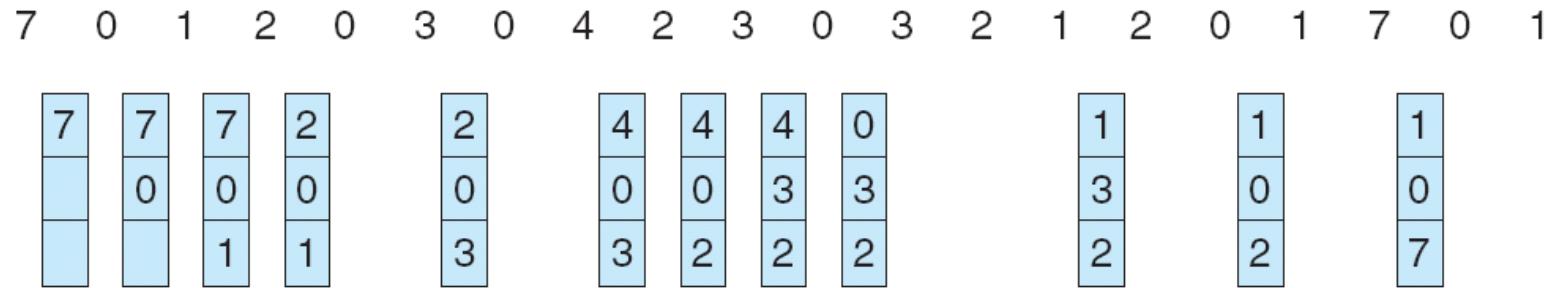




LRU Page Replacement

demo 15
optimal 9
LRU = 12

reference string



page frames

เกิด 12 page fault

- Better than FIFO but worse than Optimal Algorithm





Counting Algorithms

- Keep a counter of the number of references that have been made to each page
(on the left side)
count the frequency of usages (or references) in the part
Least frequency
- **LFU Algorithm:** replaces page with smallest count
- **MFU Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used

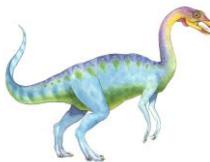
ໃນກູດເວົ້າອີງຕົວ: ເປົ້ານັ້ນ

See my note.

LFU : Least Frequently Used (ຄວາມຄືໃນການຄູກໃຫ້ນ້ອຍທີ່ສຸດ)

MFU: Most Frequently Used (ຄວາມຄືໃນການຄູກໃຫ້ມາກທີ່ສຸດ)





Allocation of Frames

- Each process needs *minimum* number of pages IBM's SS format (substring record)
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - instruction is 6 bytes, might span 2 pages 4 bytes / page
 - 2 pages to handle *from*
 - 2 pages to handle *to*
- Two major allocation schemes
 - fixed allocation defn 34
 - priority allocation defn 35





Fixed Allocation

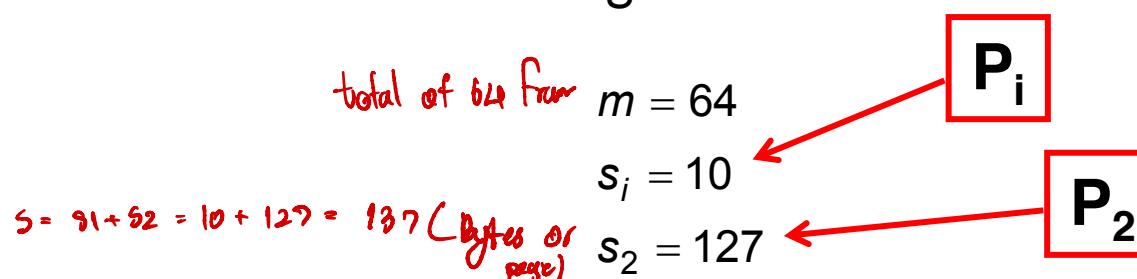
Method 1

- Equal allocation – For example, if there are 100 frames and 5 processes, give each process 20 frames. ($\text{ได้จาก (frame / process)} = (100 / 5)$)

Method 2

- Proportional allocation – Allocate according to the size of process

- s_i = size of process p_i
- $S = \sum s_i$
- m = total number of frames
- a_i = allocation for p_i = $\frac{s_i}{S} \times m$



P_1 ได้ 5 pages frame

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_1 = (s_1 / S) \times m$$

P_2 ได้ 59 pages frame

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

$$a_2 = (s_2 / S) \times m$$





Priority Allocation

Fixed allocation uses **size** for the proportionate allocation
priority allocation uses **priority** for the proportionate allocation

- Use a **proportional allocation** scheme using priorities rather than size

Higher-priority process (P_i) gets larger allocation (larger a_i) than lower-priority process (P_j) where $a_i > a_j$
Thus, the higher-priority process can be executed faster.
- If process P_i generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

The lower-priority process will be the victim to be replaced first. So starvation may occur





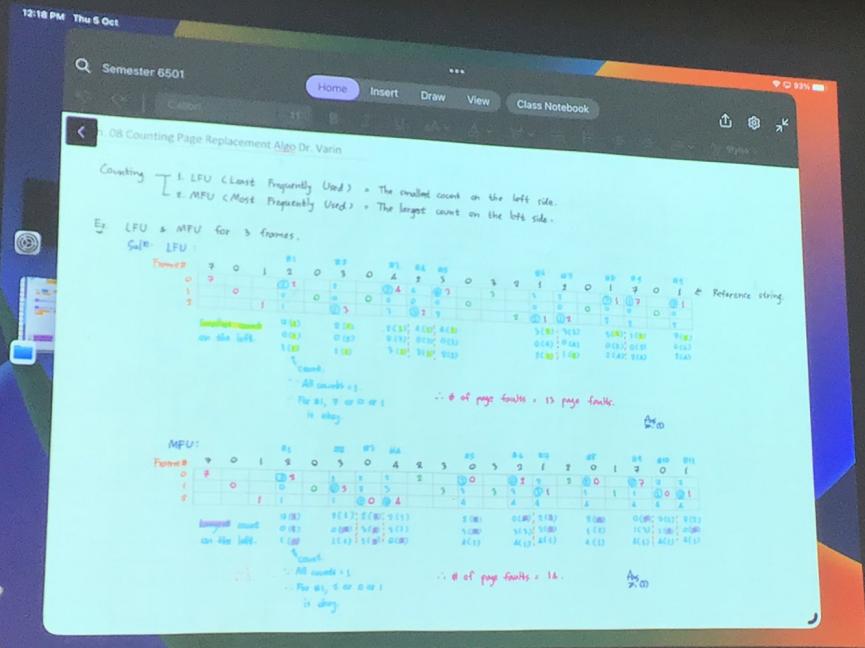
Global vs. Local Allocation

No free frame is left. There must be a replacement

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
- **Local replacement** – each process selects from only its own set of allocated frames

Select some frame of the process's own.

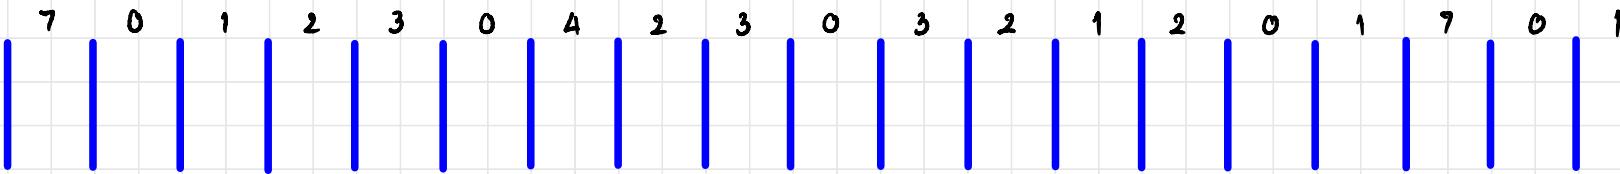




Ex LFU & MFU for 3 frames.

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

Soln LFU



End of Chapter 8

